

LECTURE 3.

SOCKET API INTRODUCTION

References:

Chapter 3 & Chapter 11: Unix network programming

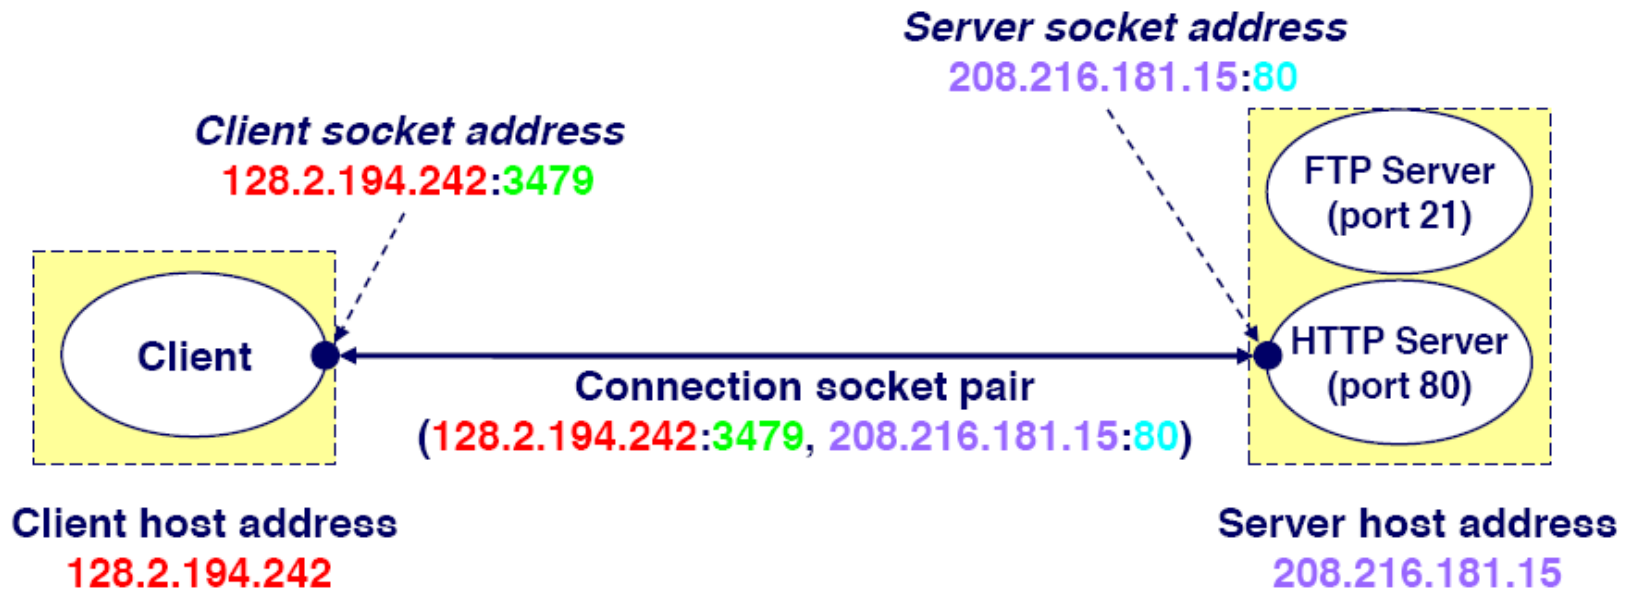
Content

- Socket
- Stream Socket
- Datagram Socket
- APIs for managing names and IP addresses
- Socket Address Structures

Socket

- What is a socket ?
- *Sockets* (in plural) are an application programming interface (API) application program at the TCP/IP stack
- A *socket* is an abstraction through which an application may send and receive data
- A socket allows an application to plug in to the network and communicate with other applications that are plugged in to the same network.

Socket (cont)



Socket (cont)

- The main types of sockets in TCP/IP are
 - *stream sockets* : use TCP as the end-to-end protocol (with IP underneath) and thus provide a reliable byte-stream service
 - *datagram sockets* : use UDP (again, with IP underneath) and thus provide a **best-effort** datagram service
- Socket Address : include host name and port

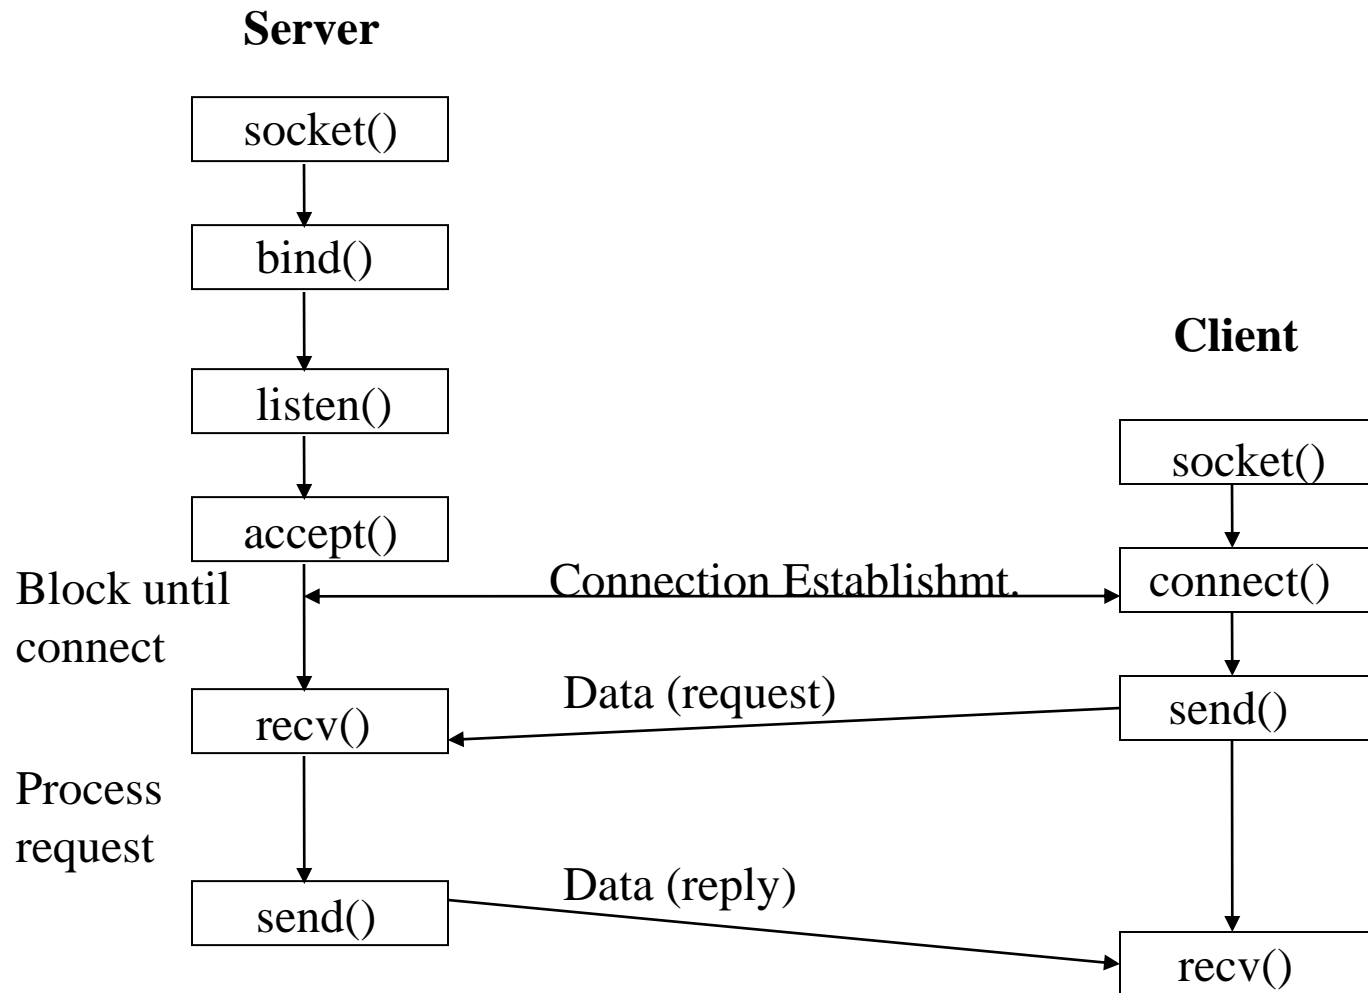
Socket: how to use

- Setup socket
 - Where is the remote machine (IP address, hostname)
 - What service gets the data (port)
- Send and Receive
 - Designed just like any other I/O in unix
 - send – write
 - recv -- read
- Close the socket

Stream sockets (TCP)

- TCP provides connections between clients and servers
- TCP also provides reliability :
 - When TCP sends data to the other end, it requires an acknowledgment in return
- TCP provides flow control
 - TCP will ensure that a sender is not overwhelming a receiver by sending packets faster than it can consume
- TCP connection is full-duplex
 - Send and receive data over single connection

Stream sockets(TCP)



Stream Socket APIs

- `socket()`
 - creates a socket of a given domain, type, protocol (buy a phone)
 - Returns a file descriptor (called a socket ID)
- `bind()`
 - Assigns a name to the socket (get a telephone number)
 - Associate a socket with an IP address and port number (Eg : 192.168.1.1:80)
- `connect()`
 - Client requests a connection request to a server
 - This is the first of the client calls

Stream Socket APIs (cont)

- `accept()` :
 - Server accept an incoming connection on a listening socket (request from a client- answer phone)
 - There are basically three styles of using `accept`:
 - *Iterating server*: only one socket is opened at a time.
 - *Forking server*: after an `accept`, a child process is forked off to handle the connection.
 - *Concurrent single server*: simultaneously wait on all open sockets, and waking up the process only when new data arrives

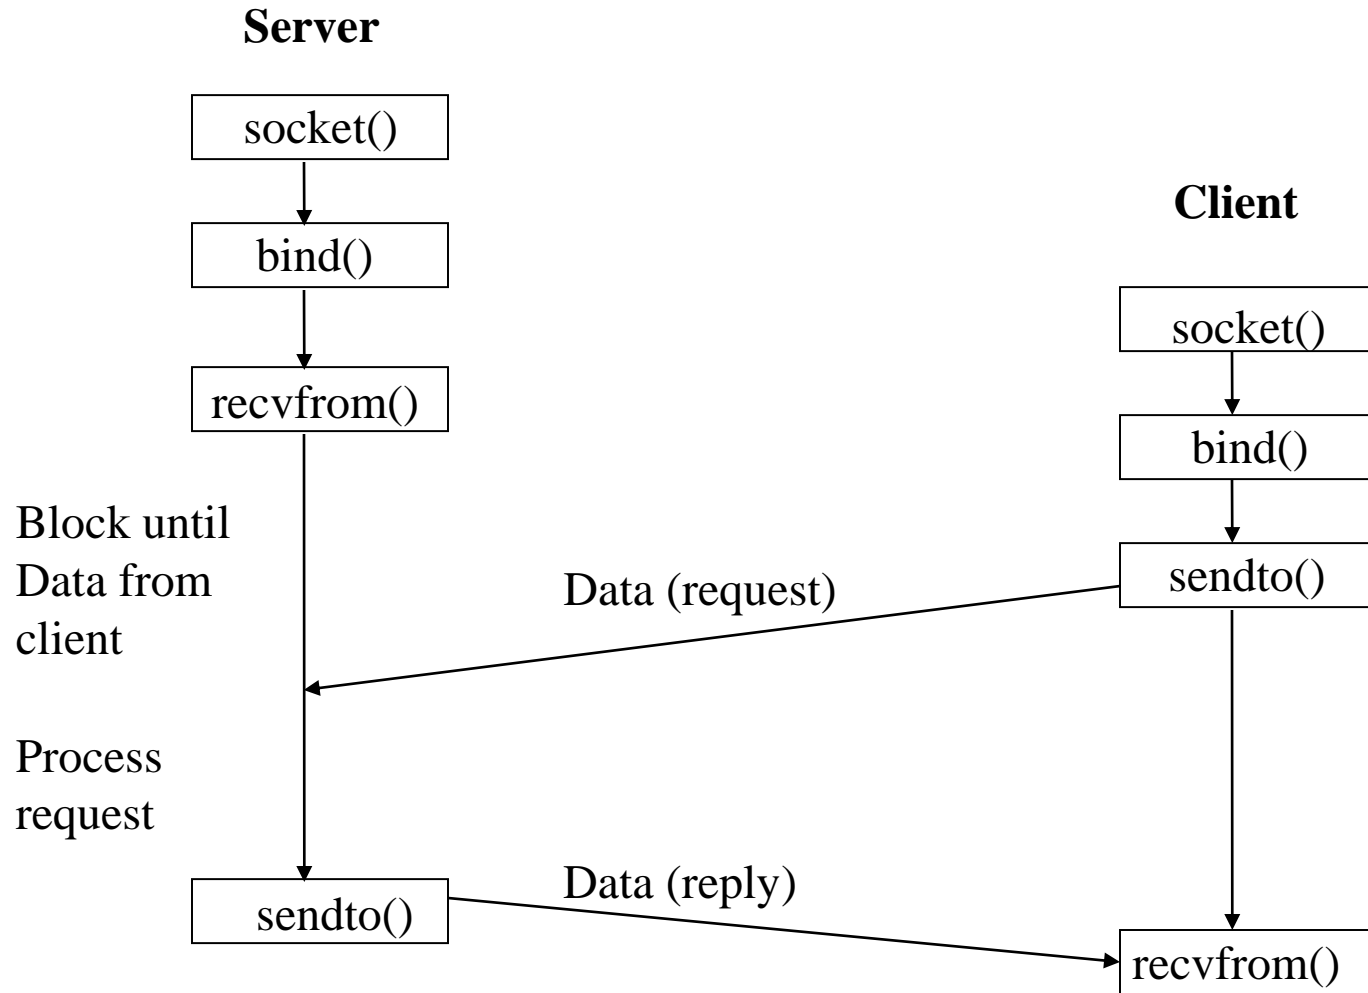
Stream Socket APIs (cont)

- `listen()`
 - Specifies the number of pending connections that can be queued for a server socket. (call waiting allowance)
- `send()`
 - Write to connection (speak)
 - Send a message
- `recv()`
 - read from connection (listen)
 - Receive data on a socket
- `close()`
 - close a socket (end the call)

Datagram Socket (UDP)

- UDP is a simple transport-layer protocol
- If a datagram is errored or lost, it won't be automatically retransmitted (can process in application)
- UDP provides a *connectionless* service, as there need not be any long-term relationship between a UDP client and server

Datagram Socket (UDP)



Socket programming in C

- `<stdio.h>`
 - input and output of basic C programs.
- `<sys/types.h>`
 - Contains definitions of data types used in system calls. These types are used in the next two include files.
- `<sys/socket.h>`
 - Includes definitions of structures needed for sockets.
- `<netinet/in.h>`
 - Contains constants and structures needed for internet domain addresses.

Socket Address Structures

- Most socket functions require a pointer to a socket address structure as an argument.
- Each supported protocol suite defines its own socket address structure.
- A Socket Address Structure is a structure which has information of a socket to create or connect with it
- Different types of socket address structures
 - IPv4
 - IPv6

IPv4 socket address structure

```
#include <netinet/in.h>
struct in_addr {
    in_addr_t s_addr;           // 32-bit IPv4 address
                                // network byte ordered
};

struct sockaddr_in {
    uint8_t sin_len;            // length of structure
    sa_family_t sin_family;     // AF_INET
    in_port_t sin_port;         // 16-bit TCP or UDP port number
                                // network byte ordered
    struct in_addr sin_addr;    // 32-bit IPv4 address
                                // network byte ordered
    char sin_zero[8];           // unused
};
```


Address families in sys/socket.h

```
/*  
 * Address families.  
 */  
#define AF_UNSPEC      0          /* unspecified */  
#define AF_LOCAL      1          /* local to host (pipes, portals) */  
#define AF_UNIX        AF_LOCAL  /* backward compatibility */  
#define AF_INET        2          /* internetwork: UDP, TCP, etc. */  
#define AF_IMPLINK     3          /* arpanet imp addresses */  
#define AF_PUP         4          /* pup protocols: e.g. BSP */  
#define AF_CHAOS       5          /* mit CHAOS protocols */  
#define AF_NS          6          /* XEROX NS protocols */  
#define AF_ISO         7          /* ISO protocols */  
#define AF_OSI         AF_ISO    /* OSI protocols */  
#define AF_ECMA        8          /* European computer manufacturers */
```

IPv4 socket address structure

```
#include <netinet/in.h>
```

```
struct in_addr {  
    in_addr_t s_addr;    /* 32-bit IPv4 address */  
                        /* network byte ordered */  
};
```

```
struct sockaddr_in {  
    uint8_t sin_len;    /* length of structure */  
    sa_family_t sin_family; /* AF_INET */  
    in_port_t sin_port; /* 16-bit TCP or UDP port number */  
                        /* network byte ordered */  
    struct in_addr sin_addr; /* 32-bit IPv4 address */  
                        /* network byte ordered */  
    char sin_zero[8]; /* unused */  
};
```

- **in_addr_t** is equivalent to the type **uint32_t**
- **uint8_t, uint16_t, uint32_t**: Integer type with a width of exactly 8, 16, 32 bits.

IPv6 socket address structure

```
#include <netinet/in.h>
struct in6_addr {
    uint8_t s6_addr[16];    // 128-bit IPv6 address
                             // network byte ordered
};

#define SIN6_LEN    // required for compile-time tests

struct sockaddr_in6 {
    uint8_t sin6_len;        // length of this struct
    sa_family_t sin6_family; // AF_INET6
    in_port_t sin6_port;     // transport layer port#
                             // network byte ordered
    uint32_t sin6_flowinfo;  // flow information, undefined
    struct in6_addr sin6_addr; // IPv6 address
                             // network byte ordered
    uint32_t sin6_scope_id;  // set of interfaces for a scope
};
```

IPv6 socket address structure

```
#include <netinet/in.h>
```

```
struct in6_addr {  
    uint8_t s6_addr[16]; /* 128-bit IPv6 address */  
                        /* network byte ordered */ };
```

```
#define SIN6_LEN      /* required for compile-time tests */
```

```
struct sockaddr_in6 {  
    uint8_t sin6_len; /* length of this struct */  
    sa_family_t sin6_family; /* AF_INET6 */  
    in_port_t sin6_port; /* transport layer port# */  
                        /* network byte ordered */  
    uint32_t sin6_flowinfo; /* flow information, undefined */  
    struct in6_addr sin6_addr; /* IPv6 address */  
                        /* network byte ordered */  
    uint32_t sin6_scope_id; /* set of interfaces for a scope */ };
```

IP Number translation

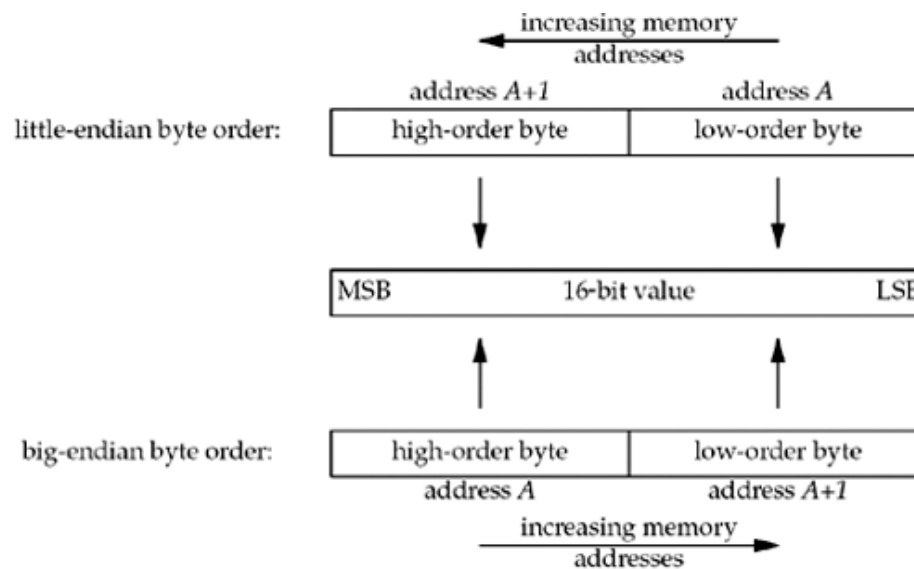
- IP address strings to 32 bit number
- Hence, these routines translate between the address as a string and the address as the number.
- Hence, we have 4 representations:
 - IP number in host order
 - IP number in network order
 - Presentation (eg. dotted decimal)
 - Fully qualified domain name

APIs for managing names and IP addresses

- Auxiliary API:
 - All binary values are network byte ordered
 - `htons()`, `htonl()`, `ntohs()`, `ntohl()`: **byte ordering**
 - `inet_ntoa()`, `inet_aton()`: Convert IPv4 addresses from a dots-and-number string (eg : 192.168.1.1) to a struct `in_addr` and back
 - `inet_pton()`, `inet_ntop()`: conversion of IPv4 or IPv6 numbers between presentation and strings

Byte Ordering

- There are two ways to store the two bytes in memory
 - little-endian byte order
 - big-endian byte order



Byte Ordering (cont)

- There is no standard between these two byte-orderings
 - The Internet protocols use big-endian byte ordering
 - Host order can be big- or little-endian
 - X86: little-endian
 - SPARC: big-endian
- Conversion
 - *htons(), htonl(): host to network short/long*
 - *ntohs(), ntohl(): network order to host short/long*
- What need to be converted?
 - Address, port?
 - Because destination host reads address, TCP/UDP port number from IP,TCP packets sent from source

htons(), htonl(), ntohs(), ntohl()

- Convert multi-byte integer types from host byte order to network byte order

```
#include <netinet/in.h>
uint32_t htonl(u_long hostlong); // host to network long
uint16_t htons(u_short hostshort); // host to network short
uint32_t ntohl(u_long netlong); // network to host long
uint16_t ntohs(u_short netshort); // network to host short
```

- Each function returns the converted value.

inet_ntoa()

```
#include <arpa/inet.h>
char *inet_ntoa(struct in_addr in);
```

- Convert IP addresses from a struct `in_addr` to a dots-and-number string
- Return: the dots-and-numbers string

```
struct in_addr someAddr;
if(inet_aton("10.0.0.1", someAddr))
    printf("The address is valid");
else printf ("The address is invalid");
char *addrStr;
addrStr = inet_ntoa(someAddr);
```

inet_aton()

```
#include <arpa/inet.h>
int inet_aton(const char *cp, struct in_addr *inp)
```

- Convert IP addresses from a dots-and-number string to a struct `in_addr`
- Return:
 - The value non-zero if the address is valid
 - The value 0 if the address is invalid

```
struct in_addr someAddr;
if(inet_aton("10.0.0.1", &someAddr))
    printf("The address is valid");
else printf ("The address is invalid");
```

inet_addr()

```
#include <arpa/inet.h>
in_addr_t inet_addr(const char *cp);
```

- Convert IP addresses from a dots-and-number string to a struct `in_addr`
- Return:
 - The value -1 if there's an error
 - The address as an `in_addr_t`

```
struct in_addr someAddr;
someAddr.s_addr = inet_addr("10.0.0.1");
```

For IPv6?

inet_pton()

```
#include <arpa/inet.h>
int inet_pton(in family, const char *cp, void *addr)
```

- Convert IP addresses from a dots-and-number string to a struct `in_addr` or `in6_addr`
- `family` is `AF_INET` or `AF_INET6`
- Return:
 - The value non-zero if the address is valid
 - The value 0 if the address is invalid

inet_ntop()

```
#include <arpa/inet.h>
const char *inet_ntop(int family, const void *addr,
                      char *cp, size_t len);
```

- Convert IP addresses from a struct `in_addr` to a dots-and-number string
- Return: the dots-and-numbers string

```
struct sockaddr_in sa;
char str[INET_ADDRSTRLEN];

// store this IP address in sa:
inet_pton(AF_INET, "192.0.2.33", &(sa.sin_addr));

// now get it back and print it
inet_ntop(AF_INET, &(sa.sin_addr), str, INET_ADDRSTRLEN);
printf("%s\n", str);
```

APIs for managing names and IP addresses

- `gethostname()` : Returns the name of the system
- `gethostbyname()` : Get an IP address for a hostname, or vice-versa

ADDRESS RESOLUTION

Content

- IPv4 and IPv6
- DNS
- Address and Name APIs

IPv4

- Developed in APRANET (1960s)
- 32-bit number
- Divided into classes that describe the portion of the address assigned to the network (netID) and the portion assigned to endpoints (hosten)
 - A : netID – 8 bit
 - B : netID – 16 bit
 - C : netID – 24 bit
 - D : use for multicast
 - E : use for experiments

IPv4 problem

- IPv4 addresses is being exhausted
 - Have to map multiple private addresses to a single public IP addresses (NATs)
 - Connect 2 PCs use private address space ?
 - NAT must be aware of the underlying protocols
- Developpe a new version of IP Address : IPv6

IPv6

- IPv6 address is 128 bits
 - To subdivide the available addresses into a hierarchy of routing domains that reflect the Internet's topology
- IPv6 address is typically expressed in 16-bit chunks displayed as hexadecimal numbers separated by colons

Example : 21DA:00D3:0000:2F3B:02AA:00FF:FE28:9C5A

or : 21DA:D3:0:2F3B:2AA:FF:FE28:9C5A

DNS (Domain Name System)


- Computers use IP Addresses to connect hosts
 - What about humans ? – IP Addresses are very complex and hard to remember (for people)
- Use name instead of IP Address → Domain Name System
- Problem of DNS
 - People use names, Computers use IP Addresses → translate between two spaces
 - Domain name system must be hierarchical (for management and maintain)
- Domain name space : divide to zones

DNS (cont)

- How to translate between domain name-IP Address and reverse ?
 - DNS Resolver
 - DNS Server
- A DNS query
 - *A non-recursive query* : DNS server provides a record for a domain for which it is authoritative itself, or it provides a partial result without querying other servers
 - *A recursive query* : DNS server will fully answer the query by querying other name servers
- DNS primarily uses User Datagram Protocol (UDP) on port number 53 to serve requests

DNS (cont)

- Address resolution mechanism
 - Local system is pre-configured with the known addresses of the root server in a file of *root hints*
 - Query one of the root servers to find the server authoritative for the next level down
 - Querying level down server for the address of a DNS server with detailed knowledge of the lower level domain until reach the DNS Server return final address



DNS (cont)

- A *Resource Record* (RR) is the basic data element in the domain name system
- All records use the common format specified in RFC 1035 (in IP networks)
- **RR (Resource record) fields**
 - NAME (variable)
 - Name of the node to which this record pertains.
 - TYPE (2)
 - Type of RR. For example, MX is type 15
 - CLASS (2)
 - Class code
 - TTL (4)
 - Unsigned time in seconds that RR stays valid
 - RDLENGTH (2)
 - Length of RDATA field
 - RDATA (variable)
 - Additional RR-specific data

List of Address and Name APIs

#include <sys/socket.h>

- **gethostbyaddr()**

- Retrieve the name(s) and address corresponding to a network address.

- **gethostname()**

- Retrieve the name of the local host.

- **gethostbyname()**

- Retrieve the name(s) and address corresponding to a host name.

- **getprotobyname()**

- Retrieve the protocol name and number corresponding to a protocol name.

- **getprotobynumber()**

- Retrieve the protocol name and number corresponding to a protocol number.

- **getservbyname()**

- Retrieve the service name and port corresponding to a service name.

- **getservbyport()**

- Retrieve the service name and port corresponding to a port.

New APIs for IPv6

- Those APIs only supports IPv4 but IPv6 will be replace IPv4 in the future, so we need APIs support IPv6
- They are
 - getaddrinfo
 - getnameinfo
- These APIs have replaced the IPv4 specific routines

gethostbyaddr ()

```
#include <netdb.h>
#include <sys/socket.h>
struct hostent *gethostbyaddr (in_addr *addr, socklen_t len,
                                int family);
```

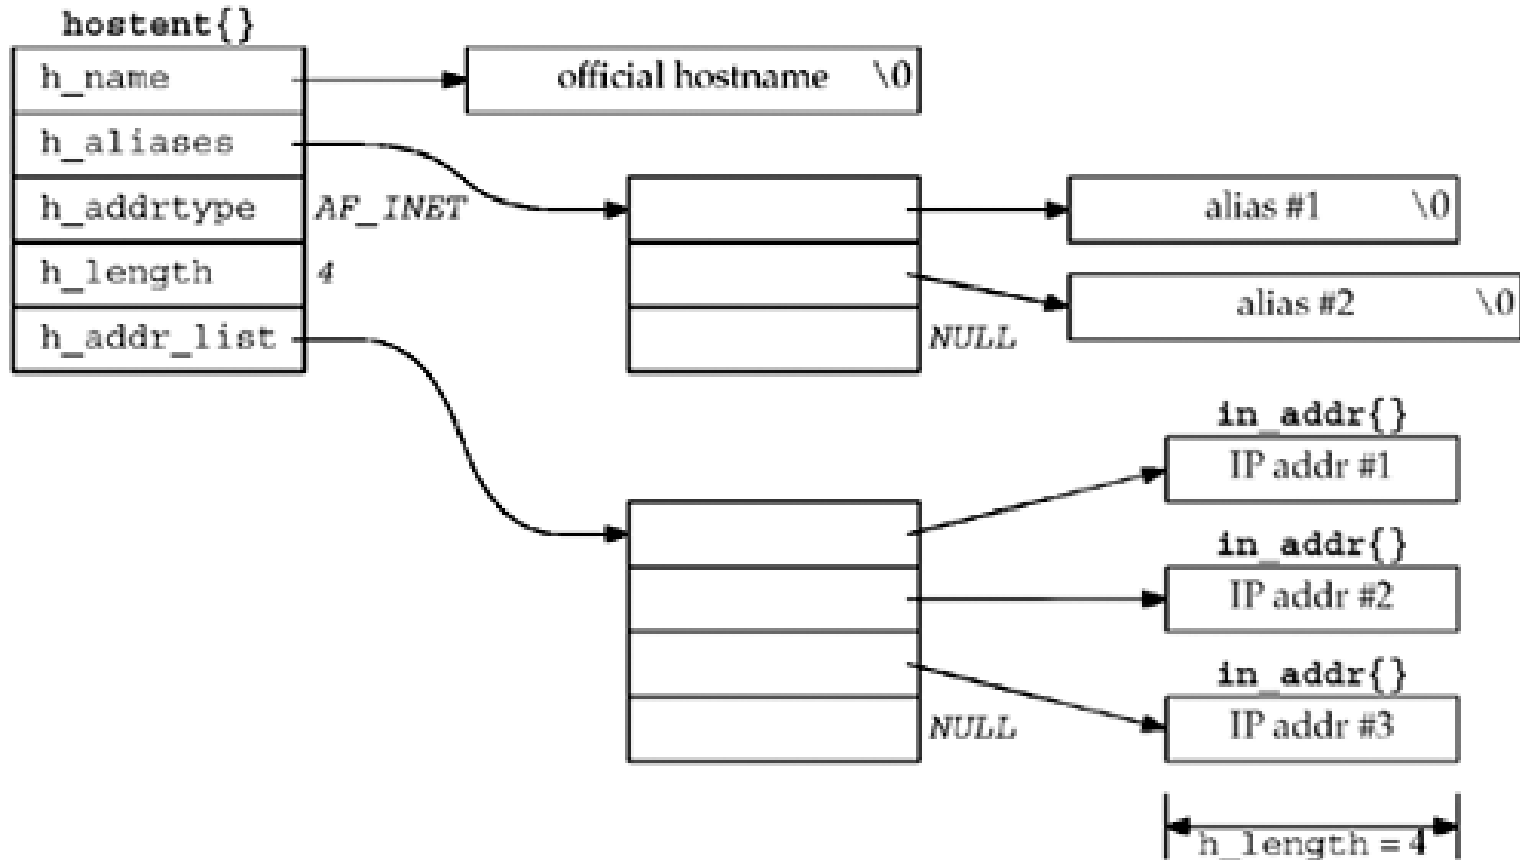
- Get host information corresponding to an address.
- Parameters:
 - [IN] `addr`: A pointer to an address in network byte order.
 - [IN] `len`: The length of the address, which must be 4 for AF_INET addresses.
 - [IN] `family`: The type of the address, which must be AF_INET.
- Return value
 - If no error occurs, returns a pointer to the `hostent` structure
 - Otherwise it returns a NULL pointer and a specific error number

struct hostent

```
struct hostent {  
    char    *h_name;        // official (canonical) name of host  
    char    **h_aliases;    // pointer to array of pointers to  
                            // alias names  
    int     h_addrtype;     // host address type: AF_INET  
    int     h_length;       // length of address: 4  
    char    **h_addr_list;  // ptr to array of ptrs with IPv4 addrs  
};
```

- What is this struct `hostent` that gets returned?
- It has a number of fields that contain information about the host in question.

struct hostent



gethostname ()

```
#include <sys/unistd.h>
#include <sys/socket.h>
int gethostname(char *name, size_t len);
```

- Return the standard host name for the local machine.
- Parameters:
 - [OUT] `name`: points to a buffer that will receive the host name.
 - [IN] `len`: the length of the buffer
- Return value
 - If no error occurs, returns 0
 - Otherwise it returns `SOCKET_ERROR` and a specific error code

gethostbyname ()

```
#include <netdb.h>
#include <sys/socket.h>
struct hostent *gethostbyname (const char *hostname);
```

- Get host information corresponding to a hostname.
- [IN] name: Points to the name of the host
- Returns a pointer to a `hostent` structure
- Return value
 - If no error occurs, returns a pointer to the `hostent` structure described above.
 - Otherwise it returns a NULL pointer and a specific error number

getservbyname()

```
#include <netdb.h>
#include <sys/socket.h>
struct servent *getservbyname (const char *servname,
                               const char *protoname);
```

- Get service information corresponding to a service name and protocol.
- Parameters:
 - [IN] `servname`: A pointer to a service name.
 - [IN] `protoname`: An optional pointer to a protocol name.
 - If this is NULL, `getservbyname()` returns the first service entry for which the name matches the `s_name` or one of the `s_aliases`.
 - Otherwise `getservbyname()` matches both the name and the proto.
- Returns
 - non-null pointer if OK
 - NULL on error

```
struct servent *sptr;
sptr = getservbyname("ftp", "tcp");
```

struct servent

```
struct servent {  
    char *s_name;  
    char **s_aliases;  
    int s_port;  
    char *s_proto;  
};
```

- s_name
 - Official name of the service.
- s_aliases
 - A NULL-terminated array of alternate names.
- s_port
 - The port number at which the service may be contacted. Port numbers are returned in network byte order.
- s_proto
 - The name of the protocol to use when contacting the service.

getservbyport ()

```
#include <netdb.h>
#include <sys/socket.h>
struct servent *getservbyport (int port, const char *protoname);
```

- Get service information corresponding to a port and protocol.
- Parameters:
 - [IN] `port`: The port for a service, in network byte order.
 - [IN] `protoname`: An optional pointer to a protocol name.
 - If this is NULL, returns the first service entry for which the port matches the `s_port`.
 - Otherwise `getservbyport()` matches both the port and the proto.
- Return
 - non-null pointer if OK
 - NULL on error

```
struct servent *sptr;
sptr = getservbyport (htons (53), "udp");
```

getpeername ()

```
#include <sys/socket.h>
int getpeername(int sockfd, struct sockaddr *addr,
                socklen_t *addr_len);
```

- Retrieve the address associated with the remote socket
- Parameters:
 - [IN] `sockfd`: the local socket connecting to remote socket
 - [OUT] `addr`: points to the `sockaddr` struct
 - [IN, OUT] `addr_len`: points to the `socklen_t` value initiated to indicate the amount of space pointed to by `addr`.
- Return:
 - On success, returns 0
 - On error, return -1 and `errno` set to indicate the error